

Ahmet Kazankaya

AI AUTOMATION ENGINEER · LLM AGENT SYSTEMS, WORKFLOW AUTOMATION, PYTHON

Istanbul, Turkey · open to relocation in Europe, remote works too · ahmetkazankaya1@gmail.com · +90 536 214 8114 · ·
linkedin.com/in/ahmetkazankaya · github.com/kasparovabi

SUMMARY

I build LLM agent systems and workflow automation that run in production, not in demos. That means Python services, data pipelines, third-party API integrations, orchestration with retries and human approval steps, and the evaluation harnesses and CI gates that decide whether an agent is allowed to ship. I own systems end to end, from the first design decision to the pager. I came to software through production work rather than a computer science degree, and my process assumes model output is wrong until something independent proves otherwise.

Since February 2026 I have consulted for a foundation running schools in 66 countries, where I own the internal automation platforms used by a communications directorate: request intake and assignment, CV screening, editorial PDF generation, and a parametric signage generator serving campuses in seven countries. Each replaced work that was previously done by hand.

AUTOMATION AND AGENT WORK

Agent orchestration. I run a fleet of scheduled agents against my own repositories and against open source projects: 141 skill definitions loaded on demand, 13 cron jobs and 6 launchd agents. A daily security audit loop has run 9 consecutive days without intervention. Agents propose; they never send. Security reports, force pushes and anything touching production data stop at a human approval gate.

Evaluation before shipping. Every agent-authored change passes a twenty-check CI gate covering correctness, security and regression before it reaches a pull request. Two of four findings my security loop produced this month were rejected by those gates before reaching a maintainer: one duplicated a published CVE, the other did not reproduce when I ran the code. That rejection rate is the argument for the gates.

Reliability work. A scheduled job failed silently for a day because the machine had no network on wake. The gap in the audit log is how I found it. I added a network precondition, a success stamp and an hourly recovery watchdog, then verified all three recovery paths rather than assuming them.

Workflow automation. Request intake that routes by learned assignment history, a CV screening pipeline in Python and Supabase, an email signature generator, a parametric signage generator producing print-ready SVG in Latin, Arabic and Cyrillic scripts, and editorial PDF generation from structured data. Each of these replaced a manual process, and each runs as a scheduled pipeline with automated checks rather than as a script someone remembers to run.

SELECTED SYSTEMS

Task management platform. Internal requests arrived by email and in conversation with no way to see who was doing what. I built intake, automatic assignment and routing that learns which colleague handles which kind of work, serving a directorate of 20 people. Next.js, React, Supabase.

Etsy factory and demand engine. A pipeline that takes a product from generation through mockup rendering to marketplace listing, with an approval gate that turns back anything failing brand or legal checks. Two marketplace API integrations, automated listing state reconciliation, and integration tests against recorded responses. Python, Etsy and Printify APIs, Telegram approval flow.

Rifki, an ESP32-S3 speech device. Hardware, firmware and the speech pipeline. Diagnosed a month-old brownout by watching the board instead of the logs: three servos drawing inrush current at boot pulled a 3.3 V rail below its 2.43 V threshold, a failure that appeared only on battery and never on USB. Deepgram, ElevenLabs, ESP-IDF, C++.

Coordinated security disclosure. Two disclosures filed under coordinated policy against open source AI tooling, the most critical rated CVSS 10.0. The vendor's own fix commit for the first cites the disclosure by name, so the outcome is verifiable in their history rather than in my write-up. A third report was prepared and closed before sending because an existing CVE already covered it.

EXPERIENCE

AI Solutions Designer, consulting · Türkiye Maarif Vakfi · Mar 2026 to now

Internal automation platforms for the Directorate of Communications.

Generative AI Engineer · Pepper Root · Nov 2025 to Jan 2026

Creative AI studio, remote.

AI Artist · Joy Game · Aug 2025 to Oct 2025

Game assets and pipeline automation.

Chief GenAI Art Director · GPT Journey · Mar 2024 to Aug 2025

AI generated commercials and custom model training for brand campaigns.

Generative AI Creator · Fiverr · Oct 2022 to Jun 2025

Graphic Designer and Photographer · Creww Digital · Jun 2020 to Nov 2022

Official Photographer · Istanbul Governorship · Jan 2018 to Apr 2020

News Editor · Skala Media Solutions · Feb 2017 to Jan 2018

STACK

Languages: Python, TypeScript, JavaScript, Swift, C++, SQL, Shell

Backend and data: FastAPI, NestJS, PostgreSQL, SQLite, Supabase, Prisma, Redis, SQLAlchemy

Frontend: Next.js, React, Tailwind, Vite, SwiftUI, Tauri, Expo

AI and agents: LLM APIs, agent orchestration, prompt evaluation, RAG, Deepgram, ElevenLabs, OpenRouter

Infrastructure: Docker, AWS EC2, Cloudflare Workers, Tailscale, launchd and cron, CI/CD, Git

Integrations: REST APIs, Microsoft Graph, Outlook 365, Etsy, Printify, Telegram, WhatsApp

Testing: pytest, vitest, XCTest, Playwright, unit and integration tests, automated checks in headless CI

EDUCATION

Mechanical Engineering, Namik Kemal University, from 2012. I left before finishing and learned the rest by building. The engineering years are where the habit of taking a system apart before judging it comes from.

Languages: Turkish native, English working proficiency.